THE UNIVERSITY OF SYDNEY

Change Management Plan

Introduction

Change management (CM) is the structured approach to moving individuals, teams and the organisation from a current state to a desired future state. It focuses on how people adopt a change, not on building the technology itself (The University of Sydney, 2026). Project management asks *"Is the system built and working?"*; change management asks *"Are people actually using it?"*. Its purpose is to make sure an IT investment delivers its benefits: a system can be deployed perfectly and still deliver nothing if people work around it, distrust it or never learn to use it.

This plan supports **Project ATLAS** (AW3, AW6), the control layer that routes each question to the most cost-effective model that meets Lyra's quality bar, caches repeat work, and measures cost and quality continuously. Students see little change, but for Lyra's staff the change is significant. Engineers must ship AI changes through evaluation gates, learning designers become part of quality assurance, and Customer Success must explain cost and quality data to universities. University academics and security officers must also be persuaded that a 'cheaper model' does not mean worse answers.

The plan combines **Kotter's 8-step model** (Kotter, 2012) to sequence the organisational change with **ADKAR** (Hiatt, 2006) to support each person from awareness to reinforcement. It is tied to the schedule (AW4), stakeholder plan (AW5) and governance framework (AW7), and change readiness is a formal criterion at Gate 3. Figure 1 and Table 1 show how the two models are sequenced and who owns each phase.

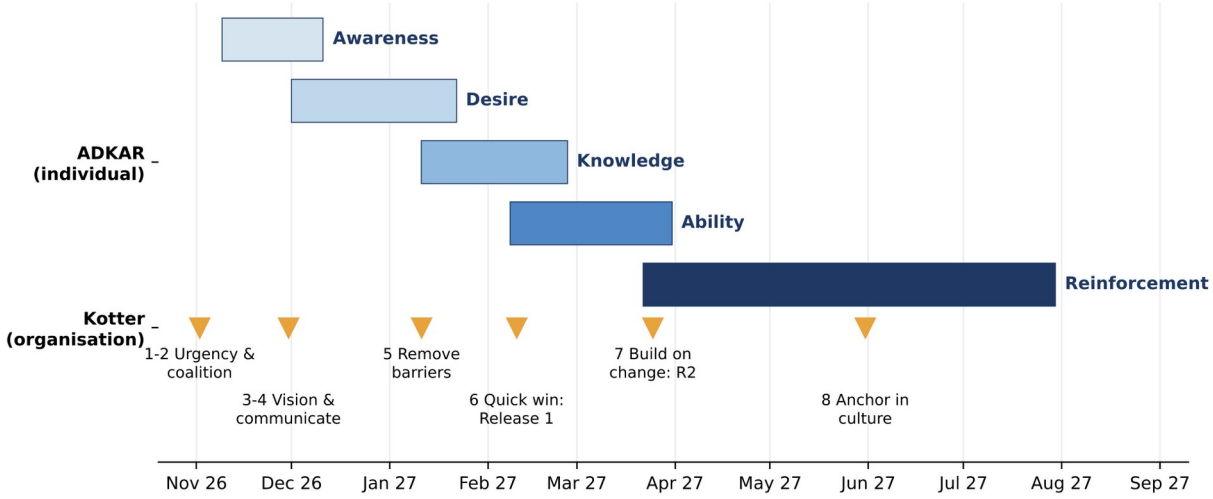


Figure 1. Change roadmap: ADKAR stages for individuals mapped against Kotter's steps for the organisation, November 2026 to July 2027

Table 1. Change roadmap by project phase (dates from AW4)

Phase and timing	Kotter steps	ADKAR focus	Key change actions	Owner
P0–P1 2 Nov – 4 Dec 2026 (Gates 0–1)	1 Urgency, 2 Coalition, 3 Vision, 4 Communicate	Awareness	Share the 44% margin gap; form champions network; CEO all-hands launch (16 Nov); co-design quality rubric	CEO and CTO (Sponsor)
P2 7 Dec 2026 – 29 Jan 2027 (Gate 2)	5 Remove barriers	Desire, Knowledge	Fast evaluation gates in CI; gate and gateway workshops; university change notices	Applied AI Lead; Engineering Manager
P3 1 Feb – 25 Mar 2027 (Gate 3, R1 10 Feb, R2 25 Mar)	6 Short-term win (R1), 7 Build on change (R2)	Ability	Game days before R1; shadow-mode evidence; canary rollout; academic webinar	Delivery Lead; Head of Customer Success
P4 and benefits 29 Mar – 30 Jul 2027 (Gate 4, review)	8 Anchor in culture	Reinforcement	Hypercare; update AI Policy and handbook; celebrate wins; post-implementation review	People & Culture Lead; Steering Committee

Change vision and objectives

> ***Vision:*** *"Every answer Lyra gives is one we can stand behind, and every dollar we spend producing it is one we can account for. ATLAS lets us grow to every Australian university without trading away the trust that brought them to us."*

The vision deliberately puts **trust first and cost second**. That framing speaks to the people most likely to resist – engineers and academics who care about quality – and reflects the 'quality before cost' principle in AW7. It is short enough for every staff member to repeat, which is Kotter's test of a clear vision (Kotter, 2012). The change management objectives are:

Table 2. Change management objectives

Objective	Target
O1 Understanding	By Gate 2 (29 January 2027), at least 90% of Lyra staff can explain why ATLAS exists and what it changes (pulse survey)
O2 Adoption of new practices	From Release 1 (10 February 2027), 100% of AI-affecting changes go through evaluation gates with no bypasses, and 100% of model calls go through the gateway
O3 Capability	Before go-live, every on-call engineer is trained and certified on ATLAS runbooks, and both learning designers are trained in annotation and rubric grading
O4 External confidence	All university security reviews are completed before Release 1, and no university raises a formal objection to Release 2
O5 Sustainment	By July 2027, the new practices are written into policy and still followed, and at least 75% of staff agree that ATLAS 'makes my work better or easier'

Change drivers and impact analysis

Six drivers make this change necessary now:

- **Financial sustainability (primary driver)** – gross margin is about 44% against a 70% target, because inference costs about AUD 0.053 per question (AW3). Growth would widen the gap.
- **Competitive advantage** – LMS vendors are adding generic AI features and free consumer tools are everywhere. Lyra wins by being *provably* grounded and affordable at campus scale.
- **Customer expectations and assurance** – universities increasingly ask vendors for evidence of AI quality and governance, as sector guidance pushes them to manage AI risk deliberately (Lodge et al., 2023).
- **Regulatory change** – automated-decision transparency obligations from 10 December 2026, the statutory privacy tort and Australia's *Guidance for AI Adoption* (National Artificial Intelligence Centre [NAIC], 2025).
- **Operational resilience** – three provider outages in Semester 1, 2026 exposed a single-provider dependency.

- **Technology opportunity** – cheap capable small models, prompt caching and Australian-hosted inference make the solution feasible only now (Anthropic, 2026).

Table 3. Change impact analysis by area

Area	Current state	Future state with ATLAS	Impact
Processes	AI changes released on reviewer judgement; manual offline evaluation before major releases; cost reviewed in monthly accounts	Every AI change passes automated evaluation gates; shadow and canary releases; weekly cost-per-question review; threshold changes approved by the AI Quality & Ethics Review Board	High
Technology	Application calls one provider directly; no tracing or caching	All calls through the LiteLLM gateway; Langfuse traces; router; semantic cache; FinOps dashboards	High for engineers, Low for users
People and roles	Engineers own features end to end; learning designers create content; Customer Success reports usage only	Engineers also own evaluation results and cost; learning designers become quality evaluators; new LLMOps on-call duties; Customer Success presents quality and cost evidence	Medium–High
Culture	'Ship it and see'; quality assumed from using the best model	'Measure, then ship'; quality and cost are joint engineering responsibilities	Medium–High
Customers (universities)	Trust based on Lyra's claims and offline evaluation	Trust based on per-tenant quality and cost evidence; formal change notices	Medium

Stakeholder analysis

Stakeholders come from the stakeholder register and power/interest map in AW5. Table 4 maps the impact on each group, its expected resistance and the ADKAR stage where it is most likely to get stuck, which shows where to focus effort. As in AW5, the highest priority groups are **unit coordinators and academics** and **university IT security and privacy officers**: both are close to the 'manage closely' quadrant and can halt adoption or renewal if they feel a change was hidden from them.

Table 4. Change impact map for affected stakeholders

Stakeholder	Impact on them	Resistance	Likely ADKAR barrier	Engagement strategy
ML and platform engineers outside the project team (about 12)	Gateway standard; evaluation gates on every AI change; new on-call runbooks	Medium	Desire: 'gates slow us down'	Co-design gate thresholds; gates under 15 minutes; show regressions caught; champion in each squad
Project team engineers	New tools (LiteLLM, Langfuse, Valkey); router ML work	Low	Knowledge: new tooling	Training budget, pairing, time-boxed spikes
Learning designers (2)	New role as evaluators; annotation workload	Low–Medium	Ability: time and tooling	Protected 20% allocation; annotation queue designed with them; visible credit
Customer Success	Must interpret and present quality and cost data	Low	Knowledge: data literacy	Dashboard training; business-review templates; FAQ
Sales	Can no longer promise 'always the best model'	Medium	Desire: losing a selling point	Reframe as 'provably grounded'; approved claims list; early reference customers
Product and Academic Advisory Board	Own the quality floor and fairness tests	Low	Awareness of new decision rights	AI Quality & Ethics Review Board terms of reference; briefing paper
Unit coordinators and academics	Students' answers may come from a smaller model	Medium–High	Awareness and Desire: distrust of 'cheap AI'	Shadow-mode evidence; academic webinar; answer-flag button; partner pilot universities
University IT security and privacy officers	Must review an architecture change	Medium	Awareness: need full disclosure	Formal change notice; updated trust pack; security briefing calls
Students	Possibly faster answers; transparency notice	Low	Awareness	Plain-language in-product notice and help article
Executives and Board	Sponsor and fund; accept governance changes	Low	Reinforcement after go-live	Gate packs, benefits dashboard, visible sponsorship

Communication plan

The change communication plan builds on the project communication plan in AW5 and focuses on the *change messages*: why, what's in it for me, and what I need to do. Two principles from the lecture apply: explain the *why*, not just the *what*, and keep channels two-way so concerns surface early (The University of Sydney, 2026). Under Kotter's model, the vision must be communicated far more often than seems necessary, and by visible leaders (Kotter, 2012).

Table 5. Change communication plan

Stakeholder	Objective	Methods	Frequency	Owner
All Lyra staff	Create urgency and share the vision (Kotter steps 1 and 4)	CEO all-hands launch; monthly demo; Notion change hub; Slack #atlas	Launch 16 Nov 2026, then monthly	CEO and CTO
Engineering (all squads)	Build desire and knowledge for gates and gateway standards	Engineering guild sessions; written standards; pairing; sprint-review demos	Fortnightly during P2–P3	Applied AI Lead; Engineering Manager
Learning designers and Academic Advisory Board	Involve them in defining quality; build ownership	Co-design workshops; annotation training; advisory board paper	Fortnightly; quarterly board	CPO
Customer Success and Sales	Tell the new value story correctly	Enablement workshop; FAQ; claims list; business-review role-play	Before R1 and R2	Head of Customer Success; Product Owner
Universities: security and privacy	Disclose the change; obtain sign-off	Formal notice letter; briefing call; trust pack	Proposed 30 days before R1 and R2, subject to contract	Head of Security, Privacy & Compliance
Universities: academics	Build confidence that quality is protected	Webinar; one-page guide; LMS announcement template; shadow-mode results	Before Semester 1; mid-semester	Customer Success with learning designers
Students	Transparency without alarm	In-product notice; help article; privacy policy update	At R2	Product Owner; Privacy Officer

Training and capacity Building

Training builds the *Knowledge* and *Ability* stages of ADKAR. Needs come from the gap between current and future roles in Table 3, validated by a short skills self-assessment in P1. Delivery mixes formal and on-the-job learning, because new practices stick best when they are applied immediately.

Table 6. Training and capacity-building plan

Audience	Training need	Delivery method	Timing	Measure of success
Project team engineers	LiteLLM configuration, Langfuse instrumentation, Valkey vector search, router training	Vendor docs and internal spikes; pair programming; 2-day LLMOps course	P1–S2	Sprint goals met; peer-review checklist passed
All engineers	Gateway usage standard; writing evaluation cases; reading gate results; canary and rollback	90-minute workshop and recorded walkthrough; runbook drills	Before Gate 2	100% completion; quiz ≥ 80%

Audience	Training need	Delivery method	Timing	Measure of success
On-call engineers	ATLAS incident runbooks; kill switch; provider failover	Game-day simulation in staging (provider outage, cache poisoning)	S4, before R1	100% certified; mitigate in drill < 15 min
Learning designers and AAB reviewers	Quality rubric, annotation in Langfuse queues, calibration	Calibration workshops with a shared answer set	P1, then monthly	Cohen's κ ≥ 0.7
Customer Success and Sales	Interpreting tenant dashboards; explaining routing honestly; approved claims	Enablement workshop, role-play, FAQ	Before R1 and R2	Role-play passed; no off-message claims
Executives and Steering Committee	Reading the FinOps and quality dashboard; their decision rights	30-minute briefing	Before Gate 2	Dashboard used in decisions

Support systems. A Notion change hub with guides, recordings and FAQs; a #atlas-help Slack channel answered within four business hours; an ATLAS champion in each squad and in Customer Success; office hours with the Applied AI Lead twice a week during P3; and four weeks of hypercare after Release 2. Champions are respected peers who answer questions and pass on feedback; this change agent network speeds adoption through peer influence (The University of Sydney, 2026). The training budget (AUD 9,000) is in AW6.

Measurement. Training is evaluated with Kirkpatrick's four levels (Kirkpatrick & Kirkpatrick, 2016): *Reaction* (post-session rating), *Learning* (workshop quiz ≥ 80% and role-play pass), *Behaviour* (evaluation gates used without bypass; all calls through the gateway) and *Results* (cost per question and quality against target).

Resistance Management

Resistance is the norm rather than the exception (The University of Sydney, 2026), and it often contains useful information, so it is treated as feedback to diagnose rather than opposition to overcome. Table 7 lists the expected sources and responses.

Table 7. Sources of resistance and response strategies

Source	Who	Why it arises	Strategy
Loss of autonomy or speed	Engineers	Gates and the CAB feel like bureaucracy in a fast start-up	Engineers set gate thresholds with the Applied AI Lead; gates run in under 15 minutes; pre-approved standard changes; celebrate regressions caught

Source	Who	Why it arises	Strategy
Quality fear ('cheap AI')	Academics, learning designers, some engineers	Belief that a smaller model means worse answers	Evidence before change: shadow-mode results; published quality floor; automatic rollback; answer-flag button
Extra workload	Learning designers, Customer Success	New tasks on top of current roles	Protected time allocations (AW4, AW6); tooling designed with them; remove lower-value tasks
Competence anxiety	Engineers new to LLMOps; Customer Success with data	Fear of looking incompetent with new tools	Psychologically safe training, pairing, game days in staging; 'no blame' incident culture
Loss of a selling point	Sales	'We always use the best model' was an easy message	New narrative: 'provably grounded at a sustainable price'; approved claims list; reference customers
Distrust of motives	Staff and universities	Suspicion that ATLAS is just cost-cutting or a step to job cuts	Lead with trust in the vision; CEO commits publicly that ATLAS is not a headcount measure; transparent benefits reporting
Procedural caution	University security teams	Any change is a risk to them; review workload	Early notice, complete trust pack, single point of contact, flexible briefing times

Two lecture models guide the timing of support. The **Kübler-Ross change curve** predicts a dip in morale and productivity when gates first block releases (e.g., "the old way worked fine"), so coaching and support are front-loaded around Release 1, and early wins such as a regression caught are shared to rebuild confidence. **Diffusion of Innovation** theory (Rogers, 2003) guides the rollout: champions and two partner universities act as early adopters, their results become social proof for the majority, and laggards get one-to-one support.

Resistance is monitored through pulse surveys, retrospective themes, champion reports, gate-bypass requests and answer-flag volumes. When a group stays stuck, ADKAR identifies the barrier: if engineers understand the gates but still bypass them, the barrier is *Desire*, not *Knowledge*, so the fix is incentives and peer influence rather than more training (Hiatt, 2006). Persistent individual resistance is handled one-to-one by the person's manager. Issues that put a milestone at risk are escalated through the AW7 matrix (Project Board at L2, Steering Committee at L3).

Cultural Considerations

Company Culture before the change

The Competing Values Framework identifies four culture types along two dimensions, flexibility versus control and internal versus external focus (Cameron & Quinn, 2011): **Clan** (collaborative, family-like), **Adhocracy** (entrepreneurial, innovative, risk-tolerant), **Market** (results-driven and competitive) and **Hierarchy** (controlled, process-driven).

Lyra today is predominantly an **Adhocracy with strong Clan elements**. It is a one-year-old start-up whose identity rests on building something new, and engineers, learning designers and former academics share cross-functional squads (Clan). Experiments ship quickly, and success has been measured by capability rather than efficiency (Adhocracy). There are few formal processes: changes are released on reviewer judgement and cost is checked monthly.

Lyra does have one evidence habit – product claims are settled by evaluation data (AW2) – but in practice it is applied only offline and occasionally. In Schein's (2010) terms, evidence is an *espoused value* that is not yet a *basic assumption*.

Desired company culture after the change

The target is a **'disciplined adhocracy'**. Adhocracy stays dominant, with a deliberate shift towards **Market** (accountability for measurable results such as cost per question and gross margin) and selected **Hierarchy** practices (evaluation gates, the CAB and audit trails). Clan values stay through shared ownership of quality between engineers and learning designers.

This target is justified for two reasons. First, Lyra's competitive advantage still depends on innovation, so a predominantly Hierarchy culture would kill the speed that lets a 34-person company outpace incumbents. Second, scaling to eight universities while keeping their trust demands that quality and cost be *measured and owned*, not assumed. A purely Market culture would also be risky, because pressure on cost could erode quality; the 'quality before cost' rule prevents this. In one sentence: **from 'ship it and see' to 'measure it, then ship it fast'.**

Bridging the gap

McKinsey's 7S model is used to align the elements that shape culture, with shared values at the centre (The University of Sydney, 2026).

Table 8. Bridging actions using the McKinsey 7S model

7S element	Bridging action
Shared values	Write 'quality before cost' into the AI Policy and engineering handbook (Kotter step 8)
Strategy and systems	Make the new way the easy way: evaluation gates automated in CI, fast and with clear feedback
Style	The CTO and Applied AI Lead put their own changes through the gates first and share results openly, including failures
Structure	Squads keep autonomy over *how* they meet quality and cost targets; only *what* must be evidenced is fixed
Staff and skills	Evaluation contributions count in performance reviews and job descriptions (People & Culture Lead)
Rituals	A weekly 'cost and quality minute' at the engineering stand-up and a shared dashboard on the office screen (Schein, 2010)

Success metrics and monitoring

Adoption is measured with **leading** indicators, which show whether people are changing, and **lagging** indicators, which show whether the change delivered results. Each KPI maps to an objective in Table 2, and all are consolidated with the project KPIs and escalation thresholds in AW12.

Table 9. Change adoption KPIs

KPI (objective)	Type	Target	Source / frequency	Owner	Escalation trigger (AW12)
Staff awareness of why ATLAS exists (O1)	Leading	≥ 90% by Gate 2	Pulse survey, monthly	People & Culture Lead	< 75% → targeted comms
Training completion: on-call and Customer Success (O3)	Leading	100% before R1	Learning records, weekly	People & Culture Lead	< 100% at Gate 3 → release held
Game-day certification of on-call engineers (O3)	Leading	100% before R1	Drill results	Engineering Manager	Not met → R1 readiness review
Calls routed through the gateway (O2)	Adoption	100% from R1	Gateway telemetry, real time	Engineering Manager	Any direct call → L2
AI changes passing evaluation gates without bypass (O2)	Adoption	100%	CI logs, weekly	Applied AI Lead	Any unlogged bypass → L2
Learning-designer agreement, Cohen's κ (O3)	Leading	≥ 0.7	Calibration, monthly	CPO	< 0.6 → recalibration
University security reviews completed before R1 (O4)	Leading	100%	Customer Success tracker, weekly	Head of Security	Outstanding 1 week before R1 → L3

KPI (objective)	Type	Target	Source / frequency	Owner	Escalation trigger (AW12)
Academic answer flags per 10,000 answers (O4)	Lagging	No increase vs R1 baseline after R2	In-product flags, weekly	Product Owner	+25% week on week → L2
Staff agreeing 'ATLAS makes my work better or easier' (O5)	Lagging	≥ 75% by Jul 2027	Pulse survey	CTO	< 60% → change plan review
Practices still followed 3 months after closure (O5)	Sustainment	Gate coverage 100%; weekly cost review held	Post-implementation review	Steering Committee	Any lapse → Steering Committee action

Feedback mechanisms. Monthly two-question pulse surveys; sprint retrospectives; the #atlas-feedback channel; fortnightly champion check-ins; the in-product answer-flag button for academics and students; and university business reviews. Each month the Delivery Lead turns the feedback into a 'you said, we did' update so people can see their input changing the plan.

Continuous improvement. The change plan follows the ITIL 4 continual improvement model used in AW12 (AXELOS, 2019), which works as a plan–do–check–act loop (Deming, 1986). At each stage gate the Steering Committee reviews adoption KPIs alongside technical readiness. Where a metric is off track, the team diagnoses the ADKAR barrier, adjusts the intervention (for example more targeted training, a revised message or a new champion) and measures again. After closure, Reinforcement continues until the July 2027 benefits review, when the change is anchored in policy and ownership passes to business-as-usual management (Kotter, 2012).

References

Anthropic. (2026). *Pricing* [Claude Platform documentation]. Retrieved October 9, 2026, from https://platform.claude.com/docs/en/about-claude/pricing

AXELOS. (2019). *ITIL foundation: ITIL 4 edition*. TSO.

Cameron, K. S., & Quinn, R. E. (2011). *Diagnosing and changing organizational culture: Based on the competing values framework* (3rd ed.). Jossey-Bass.

Deming, W. E. (1986). *Out of the crisis*. MIT Press.

Hiatt, J. M. (2006). *ADKAR: A model for change in business, government and our community*. Prosci Learning Center Publications.

Kirkpatrick, J. D., & Kirkpatrick, W. K. (2016). *Kirkpatrick's four levels of training evaluation*. ATD Press.

Kotter, J. P. (2012). *Leading change*. Harvard Business Review Press.

Lodge, J. M., Howard, S., Bearman, M., Dawson, P., & Associates. (2023). *Assessment reform for the age of artificial intelligence*. Tertiary Education Quality and Standards Agency. https://www.teqsa.gov.au/guides-resources/resources/corporate-publications/assessment-reform-age-artificial-intelligence

National Artificial Intelligence Centre. (2025). *Guidance for AI adoption*. Department of Industry, Science and Resources. https://www.industry.gov.au/publications/guidance-ai-adoption

Rogers, E. M. (2003). *Diffusion of innovations* (5th ed.). Free Press.

Schein, E. H. (2010). *Organizational culture and leadership* (4th ed.). Jossey-Bass.

The University of Sydney. (2026). *Week 07: Organisational change management* [Lecture slides]. INFO5990 Professional Practice in IT, Canvas.